

## Analytical PART:

To meet the requirements, I wrote the analysis from Step 1–6 entirely without any data mining methods, focusing only on data loading, cleaning, transformation, and Exploratory Data Analysis (EDA). After completing the exploratory phase, I applied two new data mining methods: **Random Forest Regression** and **Neural Network Regression (MLPRegressor)**.

Random Forest provided a powerful ensemble-based approach to understand feature importance and predict rental demand, while the Neural Network model captured complex non-linear relationships within the dataset. Together, these models extended the insights gained during EDA into robust predictive analytics, completing a comprehensive analytical workflow using the Bike Sharing dataset.

### **Why These Two Data Mining Methods Were Selected and How They Support the EDA Objectives:**

The Exploratory Data Analysis (EDA) revealed meaningful trends and variable relationships within the Bike Sharing (day.csv) dataset. These insights guided the selection of **Random Forest Regression** and **Neural Network Regression**, as both methods directly extend the patterns discovered during EDA into strong predictive models.

#### **Random Forest Regression- How EDA Led to This Choice**

EDA revealed several patterns:

- Rental counts increase with temperature
- Humidity and windspeed strongly affect demand
- Season and weather significantly impact user behavior
- Several variables interact in non-linear ways that may be difficult to capture using simple linear models

These observations suggested the need for a model that:

- Handles **non-linear relationships**
- Automatically captures **feature interactions**
- Provides **feature importance** to validate EDA insights

#### **How Random Forest Supports EDA Objectives**

Random Forest helps to:

- Rank which variables influence rentals the most
- Confirm EDA findings about temperature, season, and weather effects
- Produce accurate predictions using many interacting variables
- Handle outliers and non-linear patterns that EDA visualizations hinted at

It bridges the gap between visual trends (EDA) and measurable relationships (predictive modeling).

#### **Neural Network Regression- How EDA Led to This Choice**

EDA showed:

- Increasing rentals with temperature
- Sudden drops during poor weather
- Seasonal peaks and troughs
- Several variables combining in complex ways to influence demand

These observations indicated that the dataset may contain **non-linear and layered relationships** that EDA can show visually but cannot model mathematically.

A Neural Network was chosen because it:

- Learns patterns automatically from data
- Models complex relationships beyond simple linear trends
- Detects subtle interactions between temperature, weather, working days, and seasonal changes

#### **How the Neural Network Supports EDA Objectives**

The Neural Network model:

- Captures hidden patterns not obvious from charts
- Validates and extends EDA observations through predictive learning
- Learns deeper non-linear relationships suggested by EDA
- Provides an alternative modeling perspective to compare with Random Forest

#### **Python Code:**

##### **# STEP 1: IMPORT LIBRARIES**

```
import pandas as pd
```

```

import numpy as np
import matplotlib.pyplot as plt
import seaborn as sns

# ML imports
from sklearn.ensemble import RandomForestRegressor
from sklearn.model_selection import train_test_split
from sklearn.metrics import mean_squared_error, mean_absolute_error, r2_score
from sklearn.neural_network import MLPRegressor
from sklearn.preprocessing import StandardScaler

# STEP 2: LOAD day.csv FILE
uploaded = files.upload()
df = pd.read_csv("day.csv")
print("\nDataset Info:")
print(df.info())
print("\nSummary Statistics:")
print(df.describe())
print("\nMissing Values:")
print(df.isnull().sum())
print("\nDataset Shape:", df.shape)

# STEP 3: DATA CLEANING
# 1. Remove duplicate rows
duplicate_count = df.duplicated().sum()
print("\nDuplicate Rows:", duplicate_count)
if duplicate_count > 0:
    df = df.drop_duplicates()
# 2. Handle missing values using forward-fill
df = df.fillna(method='ffill')
print("\nMissing Values After Cleaning:")
print(df.isnull().sum())
# 3. Convert date column from string to datetime
df['dteday'] = pd.to_datetime(df['dteday'])
# 4. Drop irrelevant column "instant" if present
if 'instant' in df.columns:
    df.drop(columns=['instant'], inplace=True)
# 5. Rename columns for readability
df.rename(columns={'mnth': 'month', 'yr': 'year'}, inplace=True)
print("\nData Cleaning Completed.")

# STEP 4: DATA EXPLORATION
# Look at skewness and kurtosis of numeric features
print("\nSkewness:")
print(df.skew(numeric_only=True))
print("\nKurtosis:")
print(df.kurt(numeric_only=True))

# Boxplot to visually inspect outliers
plt.figure(figsize=(10,5))
sns.boxplot(data=df[['temp', 'hum', 'windspeed', 'cnt']])
plt.title("Boxplot of Key Numerical Columns")
plt.show()

# STEP 5: TRANSFORMATION
# Convert normalized temperature values into °C for human readability
df['temp_celsius'] = df['temp'] * 41
df['atemp_celsius'] = df['atemp'] * 50

# STEP 6: BASIC VISUALIZATIONS (EDA ONLY)
# 1. Distribution of rentals
plt.figure(figsize=(8,5))

```

```

sns.histplot(df['cnt'], kde=True, bins=30)
plt.title("Distribution of Bike Rentals")
plt.xlabel("Rental Count")
plt.ylabel("Frequency")
plt.show()

# 2. Correlation heatmap
plt.figure(figsize=(10,8))
sns.heatmap(df.corr(), annot=True, fmt=".2f", cmap="coolwarm", center=0)
plt.title("Correlation heatmap (numeric features)")
plt.show()

```

```

# 3. Rentals by season
plt.figure(figsize=(8,5))
sns.boxplot(x='season', y='cnt', data=df, palette='Set2')
plt.title("Bike Rentals Across Seasons")
plt.show()

```

```

# 4. Trend of rentals over time
plt.figure(figsize=(12,5))
sns.lineplot(x='dteday', y='cnt', data=df)
plt.title("Bike Rentals Over Time")
plt.show()

```

#### **# Data Mining method : RANDOM FOREST REGRESSION MODEL**

```

# Select numeric features only for modeling
numeric_df = df.select_dtypes(include=[np.number]).copy()

# X = input features, y = target
X = numeric_df.drop(columns=['cnt'])
y = numeric_df['cnt'].astype(float)

# Train-test split (80% train, 20% test)
X_train, X_test, y_train, y_test = train_test_split(
    X, y, test_size=0.2, random_state=123)

# Build Random Forest Model
rf = RandomForestRegressor(n_estimators=200, random_state=123, n_jobs=-1)
rf.fit(X_train, y_train)

# Predict on test data
rf_pred = rf.predict(X_test)

```

```

# Evaluate performance
rf_mse = mean_squared_error(y_test, rf_pred)
rf_rmse = np.sqrt(rf_mse)
rf_mae = mean_absolute_error(y_test, rf_pred)
rf_r2 = r2_score(y_test, rf_pred)
print("\nRandom Forest Regression Metrics:")
print(f" MAE : {rf_mae:.3f}")
print(f" RMSE: {rf_rmse:.3f}")
print(f" R2 : {rf_r2:.4f}")

```

```

# Feature Importance plot
importances = rf.feature_importances_
feat_names = X.columns
idx = np.argsort(importances)
plt.figure(figsize=(8,6))
plt.barh(feat_names[idx], importances[idx])
plt.grid(axis='x', linestyle='--', alpha=0.6)

```

```

plt.title("Random Forest Feature Importances")
plt.xlabel("Importance")
plt.show()
# Actual vs Predicted scatter plot
plt.figure(figsize=(8,6))
plt.scatter(y_test, rf_pred, alpha=0.7)
plt.plot([y_test.min(), y_test.max()], [y_test.min(), y_test.max()], 'r--')
plt.xlabel("Actual cnt")
plt.ylabel("Predicted cnt")
plt.title("Random Forest: Actual vs Predicted")
plt.show()

# Prediction error distribution
errors = y_test - rf_pred
plt.figure(figsize=(8,5))
sns.histplot(errors, kde=True)
plt.title("Random Forest Prediction Errors")
plt.xlabel("Error (actual - predicted)")
plt.show()

# Data Mining method : NEURAL NETWORK REGRESSION (MLPRegressor)

# Scale features for neural network
scaler = StandardScaler()
X_train_scaled = scaler.fit_transform(X_train)
X_test_scaled = scaler.transform(X_test)

# Build MLPRegressor model
mlp = MLPRegressor(hidden_layer_sizes=(64,32), max_iter=500, random_state=123)
mlp.fit(X_train_scaled, y_train)

# Predict
mlp_pred = mlp.predict(X_test_scaled)

# Evaluate
mlp_mse = mean_squared_error(y_test, mlp_pred)
mlp_rmse = np.sqrt(mlp_mse)
mlp_mae = mean_absolute_error(y_test, mlp_pred)
mlp_r2 = r2_score(y_test, mlp_pred)
print("\nMLPRegressor Metrics:")
print(f" MAE : {mlp_mae:.3f}")
print(f" RMSE: {mlp_rmse:.3f}")
print(f" R2 : {mlp_r2:.4f}")

# Actual vs predicted
plt.figure(figsize=(8,6))
plt.scatter(y_test, mlp_pred, alpha=0.7)
plt.plot([y_test.min(), y_test.max()], [y_test.min(), y_test.max()], 'r--')
plt.xlabel("Actual cnt")
plt.ylabel("Predicted cnt")
plt.title("MLPRegressor: Actual vs Predicted")
plt.show()

# Error distribution
errors_mlp = y_test - mlp_pred
plt.figure(figsize=(8,5))
sns.histplot(errors_mlp, kde=True)
plt.title("MLPRegressor Prediction Errors")
plt.xlabel("Error (actual - predicted)")
plt.show()

```

```
print("\n\n Random Forest and Neural Network Regression Completed Successfully!")
```

### OutPut:

#### Dataset Info:

```
RangeIndex: 731 entries, 0 to 730
Data columns (total 16 columns):
 #   Column   Non-Null Count Dtype  
 --- 
 0   instant   731 non-null   int64  
 1   dteday    731 non-null   object 
 2   season    731 non-null   int64  
 3   yr        731 non-null   int64  
 4   mnth      731 non-null   int64  
 5   holiday   731 non-null   int64  
 6   weekday   731 non-null   int64  
 7   workingday 731 non-null   int64  
 8   weathersit 731 non-null   int64  
 9   temp       731 non-null   float64 
 10  atemp      731 non-null   float64 
 11  hum        731 non-null   float64 
 12  windspeed  731 non-null   float64 
 13  casual     731 non-null   int64  
 14  registered 731 non-null   int64  
 15  cnt        731 non-null   int64  
 dtypes: float64(4), int64(11), object(1)
memory usage: 91.5+ KB
```

None

#### Summary Statistics:

```
instant   season   yr      mnth   holiday   weekday \
count    731.000000 731.000000 731.000000 731.000000 731.000000
mean     366.000000 2.496580 0.500684 6.519836 0.028728 2.997264
std      211.165812 1.110807 0.500342 3.451913 0.167155 2.004787
min      1.000000 1.000000 0.000000 1.000000 0.000000 0.000000
25%     183.500000 2.000000 0.000000 4.000000 0.000000 1.000000
50%     366.000000 3.000000 1.000000 7.000000 0.000000 3.000000
75%     548.500000 3.000000 1.000000 10.000000 0.000000 5.000000
max     731.000000 4.000000 1.000000 12.000000 1.000000 6.000000
```

```
workingday   weathersit   temp   atemp   hum   windspeed \
count    731.000000 731.000000 731.000000 731.000000 731.000000 731.000000
mean     0.683995 1.395349 0.495385 0.474354 0.627894 0.190486
std      0.465233 0.544894 0.183051 0.162961 0.142429 0.077498
min      0.000000 1.000000 0.059130 0.079070 0.000000 0.022392
25%     0.000000 1.000000 0.337083 0.337842 0.520000 0.134950
50%     1.000000 1.000000 0.498333 0.486733 0.626667 0.180975
75%     1.000000 2.000000 0.655417 0.608602 0.730209 0.233214
max     1.000000 3.000000 0.861667 0.840896 0.972500 0.507463
```

```
casual   registered   cnt
count    731.000000 731.000000 731.000000
mean     848.176471 3656.172367 4504.348837
std      686.622488 1560.256377 1937.211452
min      2.000000 20.000000 22.000000
25%     315.500000 2497.000000 3152.000000
50%     713.000000 3662.000000 4548.000000
75%     1096.000000 4776.500000 5956.000000
max     3410.000000 6946.000000 8714.000000
```

#### Missing Values:

```
instant   0
dteday    0
season   0
yr        0
mnth     0
holiday   0
weekday   0
workingday 0
weathersit 0
temp      0
```

```
atemp      0  
hum        0  
windspeed  0  
casual     0  
registered 0  
cnt        0  
dtype: int64
```

Dataset Shape: (731, 16)

Duplicate Rows: 0

Missing Values After Cleaning:

```
instant    0  
dteday     0  
season     0  
yr         0  
mnth       0  
holiday    0  
weekday   0  
workingday 0  
weathersit 0  
temp       0  
atemp      0  
hum        0  
windspeed  0  
casual     0  
registered 0  
cnt        0  
dtype: int64
```



Data Cleaning Completed.



Skewness:


```
season    -0.000384  
year      -0.002742  
month     -0.008149  
holiday   5.654224  
weekday   0.002742  
workingday -0.793147  
weathersit 0.957385  
temp      -0.054521  
atemp     -0.131088  
hum       -0.069783  
windspeed 0.677345  
casual    1.266454  
registered 0.043659  
cnt       -0.047353  
dtype: float64
```

Kurtosis:

```
season    -1.342601  
year      -2.005487  
month     -1.209112  
holiday   30.052462  
weekday   -1.254282  
workingday -1.374686  
weathersit -0.136467  
temp      -1.118864  
atemp     -0.985131  
hum       -0.064530  
windspeed 0.410922  
casual    1.322074  
registered -0.713097  
cnt       -0.811922  
dtype: float64
```








#### Random Forest Regression

##### **Random Forest Regression Metrics:**


MAE : 525.838

RMSE: 711.411


R2 : 0.8517



Random Forest: Actual vs Predicted



Random Forest Prediction Errors





===== Neural Network Regression (MLPRegressor) =====

MLPRegressor Metrics:

MAE : 959.643

RMSE: 1230.695

R2 : 0.5561



Random Forest and Neural Network Regression Completed Successfully!